



Trinity College Dublin

Coláiste na Tríonóide, Baile Átha Cliath

The University of Dublin

School of Computer Science and Statistics

CS7CS4 Machine Learning Week 1 Assignment

Boris Flesch

20300025

October 8, 2020

MSc Computer Science, Intelligent Systems

1 Downloaded Dataset

id:2-1018.8-10

2 Questions

- (a) All sections of the Python code have been commented in accordance with each question of the assignment (see appendix A.1 or attached ZIP file). Figure 1 shows a visual representation of the data read from the downloaded CSV file:




Figure 1

- (b) (i) The learning rate values 0.001, 0.01, 0.1 and 1 all seem to be converging but at a different speed (Figure 2). More iterations would be required to observe the evolution of the cost function for $\alpha = 0.01$ and $\alpha = 0.001$.
Note about θ initial value: after having executed several tests with different initial values for θ_0 and θ_1 , that seems to have an insignificant impact on the values of the trained model parameters. The only observation that I have had is that a big initial value (i.e. $\theta_0 = \theta_1 = 1000$) will prevent the gradient descent from converging with a learning rate $\alpha = 1$. Therefore, I decided to keep initial values $\theta_0 = \theta_1 = 0$.

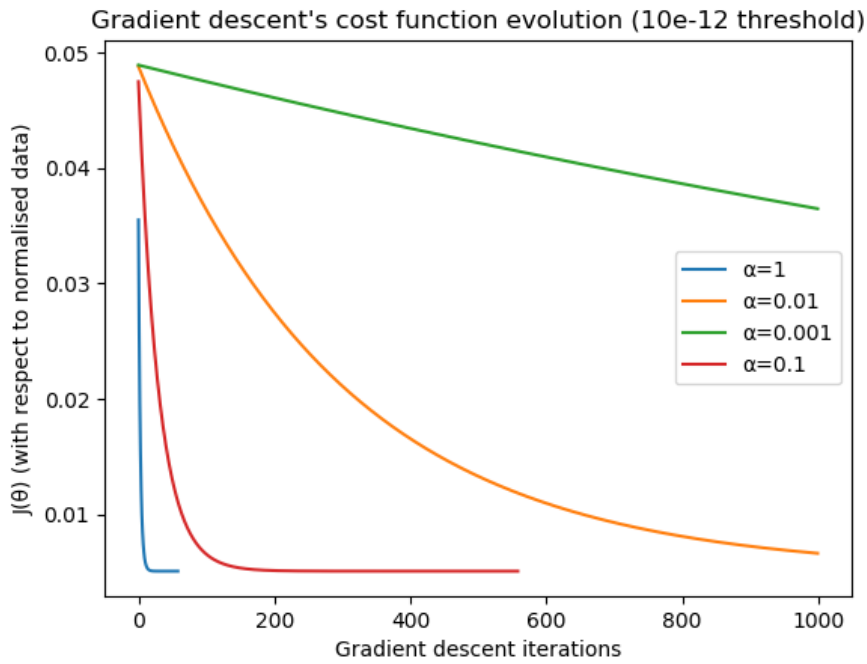


Figure 2

- (ii) After having been trained on my downloaded data, the parameter values of my linear regression model are the following: $\theta_0 = 518.027317$, $\theta_1 = -4.964503$
- (iii) The value of the cost function of my linear regression model is $J(\theta) = 172100.778305$. As a baseline model, I chose $\theta_0 = \text{"mean of output y values"}$ and $\theta_1 = 0$. This baseline model gives $J(\theta_{baseline}) = 1652762.172120$.

As we can see on Figure 3, $J(\theta_{baseline}) \gg J(\theta)$. This confirms that the linear regression model is much more accurate than the baseline model which does actually not follow the trend/path of the training data (see green dashed line on Figure 3).




Figure 3

- (iv) The linear regression model trained using sklearn (Figure 4) has the following parameter and cost function values: $\theta_0 = 518.198839$, $\theta_1 = -4.964907$ and $J(\theta) = 172100.768521$. These values are very close to the ones obtained using the gradient descent algorithm in vanilla Python, thereby confirming the accuracy of the linear regression model using gradient descent.

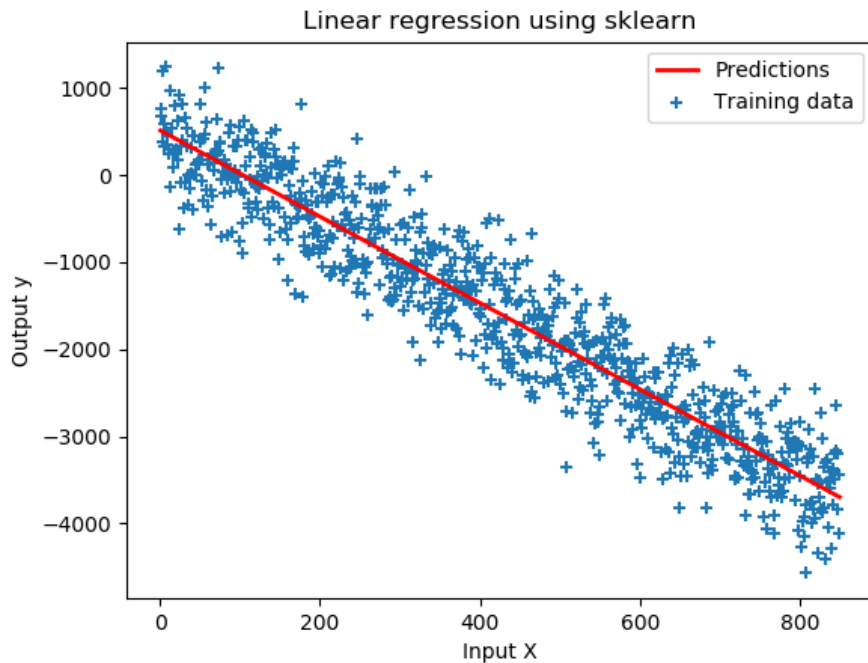


Figure 4

A Appendix

A.1 Python Code

```

1  #!/usr/bin/env python
2  # coding: utf-8
3
4  # CS7CS4/CSU44061 Machine Learning
5  # Week 1 Assignment
6  # Boris FLesch (20300025)
7  #
8  # Downloaded dataset
9  # id:2-1018.8--10
10
11 import numpy as np
12 import pandas as pd
13 import matplotlib.pyplot as plt
14
15 # (a)(i) Read data
16 def readData(file):
17     """
18     Read data from a file and outputs reshaped arrays
19
20     :param file: path to the source file
21     :return: 2 numpy arrays X and y
22     """
23     df = pd.read_csv(file, comment="#")
24     X = np.array(df.iloc[:, 0]).reshape(-1,1)
25     y = np.array(df.iloc[:, 1]).reshape(-1,1)
26
27     return X, y
28
29 X, y = readData("week1.csv")
30
31 # Visualise data
32 plt.plot(X, y, '+')
33 plt.title("Original dataset visualisation")
34 plt.xlabel("X")
35 plt.ylabel("y")
36 plt.show()
37 #plt.savefig("original-dataset.png")
38
39
40 # (a)(ii) Normalise data
41 def normaliseData(X):
42     """
43     Normalise data using average as shift and max-min as scaling factor
44

```

```

45     :param X: Data array to normalise
46     :return: Normalised data, shift value, scaling factor value
47     """
48     shift = np.average(X)
49     scalingFactor = np.max(X) - np.min(X)
50     X = (X - shift) / scalingFactor
51     return X, shift, scalingFactor
52
53 # Linear Regression
54 def linearRegression(X_in, y_in, alpha=[0.1], theta=[0,0],
55 ↪ costThreshold=10e-10, maxIterations=1000):
56     """
57     Process a linear regression on data
58
59     :param X_in: X data (array)
60     :param y_in: y data (array)
61     :param alpha: Value of the learning rate
62     :param theta: Initial values for theta (array)
63     :param costThreshold: Cost function limit between two iterations of
64 ↪ the gradient descent
65     :param maxIterations: Maximum number of iterations for the gradient
66 ↪ descent
67     :return: theta (mapped to original data, i.e. "de-normalised"),
68 ↪ errors (with respect to normalised data), finalError (mapped to
69 ↪ original data)
70     """
71
72     # Copy of input arrays
73     X = np.copy(X_in)
74     y = np.copy(y_in)
75
76     # Normalise data
77     X, shiftX, scalingFactorX = normaliseData(X)
78     y, shiftY, scalingFactorY = normaliseData(y)
79
80     # (a)(iii) Gradient Descent
81     costs = []
82     m = X.size
83
84     for k in range(maxIterations):
85         # Theta calculation
86         theta[0] += (-2*alpha / m) * sum((theta[0] + theta[1] * X) - y)
87         theta[1] += (-2*alpha / m) * sum((theta[0] + theta[1] * X) - y)
88         ↪ * X)
89
90         # Cost calculation
91         cost = 1/m * sum(np.power((theta[0] + theta[1] * X) - y, 2))
92         costs.append(cost)

```

```

87
88     if (k > 2 and abs(costs[-1] - costs[-2]) <= costThreshold):
89         break
90
91     costs = np.array(costs)
92
93     # Theta calculation (mapping value to original data)
94     theta = [
95         theta[0] * scalingFactorY + shiftY - theta[1] * shiftX *
96         ↪ scalingFactorY / scalingFactorX,
97         theta[1] / scalingFactorX * scalingFactorY
98     ]
99
100    # Final cost value (mapping value to original data)
101    finalCost = 1/m * sum(np.power((theta[0] + theta[1] * X_in) - y_in,
102    ↪ 2))
103
104    return theta, costs, finalCost
105
106    # (b)(i) Linear regression using different learning rates
107    alphas = [1, 0.01, 0.001, 0.1]
108    for alpha in alphas:
109        theta, costs, finalCost = linearRegression(X, y, alpha, theta=[0,0],
110        ↪ costThreshold=10e-12, maxIterations=1000)
111        plt.plot(range(costs.size), costs, '-', label="=" + str(alpha))
112
113    plt.legend()
114    plt.title("Gradient descent's cost function evolution (10e-12
115    ↪ threshold)")
116    plt.xlabel("Gradient descent iterations")
117    plt.ylabel("J(theta) (with respect to normalised data)")
118    plt.show()
119    #plt.savefig("cost-function-evolution.png")
120
121    # (b)(ii)
122    print("Linear regression model using vanilla Python:")
123    print("theta0 = %f ; theta1 = %f" % (theta[0], theta[1]))
124
125    # (b)(iii)
126    print("J(theta) = %f" % finalCost)
127
128    # Baseline model (theta0 = y mean, theta1 = 0)
129    bTheta = [np.mean(y), 0]
130    bCost = 1/X.size * sum(np.power((bTheta[0] + bTheta[1] * X) - y, 2))
131    print("J(theta_baseline) = %f\n" % bCost)
132
133    # Plot predictions, training data and baseline model

```

```
131 plt.title("Linear regression")
132 plt.xlabel("Input X")
133 plt.ylabel("Output y")
134 plt.scatter(X, y, marker='+', label="Training data")
135 plt.plot(X, theta[0] + theta[1] * X, 'r-', linewidth=2,
    ↪ label="Predictions")
136 plt.plot(X, bTheta[0] + bTheta[1] * X, color='g', linestyle='dashed',
    ↪ linewidth=2, label="Baseline model (y mean)")
137 plt.legend()
138 plt.show()
139 #plt.savefig("linear-regression.png")
140
141
142 # (b)(iv) Linear Regression model with sklearn
143 from sklearn.linear_model import LinearRegression
144
145 X, y = readData("week1.csv")
146
147 model = LinearRegression().fit(X, y)
148 cost = 1/X.size * sum(np.power((model.intercept_ + model.coef_ * X) - y,
    ↪ 2))
149
150 print("Linear regression model using sklearn:")
151 print("theta0 = %f ; theta1 = %f" % (model.intercept_, model.coef_))
152 print("J(theta) = %f" % cost)
153
154 plt.title("Linear regression using sklearn")
155 plt.scatter(X, y, marker='+', label="Training data")
156 plt.plot(X, model.intercept_ + model.coef_ * X, 'r-', linewidth=2,
    ↪ label="Predictions")
157 plt.legend()
158 plt.xlabel("Input X")
159 plt.ylabel("Output y")
160 plt.show()
161 #plt.savefig("linear-regression-sklearn.png")
```